

# Deliberate Attack

## Analyze attack trends

Table 1: We run a logistic model regressing success against perturb-target distance (relative to image width/height) and perturb box length (relative to image width/height) in the deliberate attack experiment. Longer perturb box length (relative to image width/height) or shorter perturb-target distance (relative to image width/height) cause success rates to significantly increase for all model and attack combinations, except for perturb box length (relative to image width/height) in untargeted attack on Cascade R-CNN. The interaction terms, even when significant, are negligibly close to 0. Table headers are explained in Appendix ??.

Group		Regression						
Attack	term	sig	estimate	std.error	statistic	p.value	conf.low	conf.high
<b>YOLOv3</b>								
Vanishing	distance	*	-6.969	1.182	-5.898	0.000	-9.308	-4.674
	length	*	7.744	0.453	17.110	0.000	6.873	8.648
	distance * length	*	-15.800	3.262	-4.844	0.000	-22.195	-9.400
Mislabeling	distance	*	-7.516	1.198	-6.276	0.000	-9.891	-5.194
	length	*	5.962	0.354	16.853	0.000	5.277	6.664
	distance * length		-3.757	2.981	-1.260	0.208	-9.565	2.127
Untargeted	distance	*	-9.214	1.496	-6.157	0.000	-12.205	-6.335
	length	*	2.276	0.257	8.860	0.000	1.775	2.782
	distance * length		3.324	2.896	1.148	0.251	-2.311	9.048
<b>SSD</b>								
Vanishing	distance	*	-11.391	1.488	-7.656	0.000	-14.362	-8.527
	length	*	4.449	0.281	15.841	0.000	3.903	5.004
	distance * length		-2.946	2.946	-1.000	0.317	-8.682	2.871
Mislabeling	distance	*	-9.002	1.596	-5.641	0.000	-12.194	-5.934
	length	*	5.846	0.303	19.286	0.000	5.257	6.446
	distance * length	*	-13.088	3.183	-4.111	0.000	-19.300	-6.813
Untargeted	distance	*	-10.046	1.760	-5.710	0.000	-13.580	-6.679
	length	*	3.873	0.280	13.823	0.000	3.328	4.426
	distance * length	*	-9.039	3.479	-2.598	0.009	-15.823	-2.178
<b>RetinaNet</b>								
Vanishing	distance	*	-22.822	2.997	-7.615	0.000	-28.866	-17.129
	length	*	2.657	0.333	7.978	0.000	2.008	3.313
	distance * length		-10.710	6.207	-1.725	0.084	-22.836	1.492
Mislabeling	distance	*	-23.172	4.211	-5.503	0.000	-31.759	-15.284

Untargeted	length	*	1.352	0.411	3.293	0.001	0.552	2.162
	distance * length		3.035	8.178	0.371	0.711	-12.833	19.200
	distance	*	-23.262	2.438	-9.540	0.000	-28.144	-18.587
	length	*	2.805	0.310	9.056	0.000	2.203	3.417
	distance * length	*	36.482	4.181	8.726	0.000	28.422	44.811
<b>Faster R-CNN</b>								
Vanishing	distance	*	-31.003	3.728	-8.317	0.000	-38.489	-23.891
	length	*	2.099	0.353	5.954	0.000	1.413	2.795
	distance * length		-0.776	7.406	-0.105	0.917	-15.251	13.769
Mislabeling	distance	*	-27.655	4.100	-6.745	0.000	-35.955	-19.909
	length	*	1.000	0.384	2.606	0.009	0.253	1.757
	distance * length		9.942	7.822	1.271	0.204	-5.291	25.353
Untargeted	distance	*	-29.452	2.752	-10.702	0.000	-34.974	-24.189
	length	*	1.949	0.294	6.625	0.000	1.375	2.529
	distance * length	*	34.982	4.521	7.738	0.000	26.255	43.974
<b>Cascade R-CNN</b>								
Vanishing	distance	*	-30.415	3.888	-7.823	0.000	-38.197	-22.974
	length	*	3.974	0.397	10.014	0.000	3.205	4.761
	distance * length	*	-25.436	8.343	-3.049	0.002	-41.832	-9.131
Mislabeling	distance	*	-34.967	4.698	-7.442	0.000	-44.420	-26.025
	length	*	1.927	0.388	4.971	0.000	1.173	2.693
	distance * length		2.023	9.140	0.221	0.825	-15.793	20.012
Untargeted	distance	*	-43.934	3.758	-11.692	0.000	-51.456	-36.729
	length	*	0.830	0.313	2.654	0.008	0.220	1.447
	distance * length	*	44.859	6.141	7.305	0.000	32.936	57.013

Table 2: We combined the data in the randomized and deliberate attack experiments to run a logistic model regressing success against object (versus non-object), with perturb-target distance (relative to image width/height) and perturb box size (relative to image width/height) as covariates. The “object” term codes object as 1 and non-object as 0. Perturbing an object (in the randomized attack) rather than a non-object (in the deliberate attack) significantly decreases success rates for all model and attack combinations, after controlling for perturb sizes and perturb-target distances. Table headers are explained in Appendix ??.

Group		Regression						
Attack	term	sig	estimate	std.error	statistic	p.value	conf.low	conf.high
<b>YOLOv3</b>								
Vanishing	object	*	-0.702	0.065	-10.865	0.000	-0.828	-0.575
	distance	*	-8.962	0.501	-17.896	0.000	-9.958	-7.995
	size	*	14.550	0.738	19.702	0.000	13.134	16.030
	distance * size	*	-38.136	4.331	-8.806	0.000	-46.722	-29.740

Mislabeling	object	*	-0.652	0.062	-10.521	0.000	-0.774	-0.531
	distance	*	-8.427	0.477	-17.661	0.000	-9.377	-7.507
	size	*	8.731	0.462	18.883	0.000	7.841	9.654
	distance * size	*	-11.042	3.119	-3.540	0.000	-17.177	-4.945
Untargeted	object	*	-0.786	0.078	-10.026	0.000	-0.940	-0.633
	distance	*	-11.538	0.789	-14.616	0.000	-13.118	-10.023
	size	*	2.107	0.265	7.957	0.000	1.589	2.627
	distance * size	*	12.922	2.700	4.786	0.000	7.649	18.235
<b>SSD</b>								
Vanishing	object		0.079	0.063	1.242	0.214	-0.045	0.203
	distance	*	-15.053	0.713	-21.121	0.000	-16.475	-13.681
	size	*	5.106	0.307	16.639	0.000	4.510	5.714
	distance * size		3.280	2.731	1.201	0.230	-2.073	8.638
Mislabeling	object	*	-0.280	0.066	-4.241	0.000	-0.409	-0.150
	distance	*	-14.939	0.800	-18.674	0.000	-16.538	-13.402
	size	*	5.517	0.310	17.805	0.000	4.915	6.130
	distance * size		-2.191	2.933	-0.747	0.455	-7.946	3.555
Untargeted	object	*	-0.220	0.070	-3.163	0.002	-0.357	-0.084
	distance	*	-15.788	0.918	-17.194	0.000	-17.628	-14.029
	size	*	3.430	0.275	12.493	0.000	2.895	3.972
	distance * size		3.110	3.178	0.979	0.328	-3.137	9.327
<b>RetinaNet</b>								
Vanishing	object	*	-0.552	0.087	-6.327	0.000	-0.724	-0.382
	distance	*	-28.013	1.755	-15.965	0.000	-31.533	-24.656
	size	*	2.736	0.339	8.081	0.000	2.078	3.406
	distance * size		3.445	5.929	0.581	0.561	-8.297	14.956
Mislabeling	object	*	-0.469	0.125	-3.748	0.000	-0.718	-0.227
	distance	*	-30.658	2.722	-11.265	0.000	-36.164	-25.501
	size	*	1.027	0.423	2.426	0.015	0.200	1.859
	distance * size	*	20.039	8.456	2.370	0.018	3.311	36.476
Untargeted	object	*	-0.359	0.082	-4.382	0.000	-0.520	-0.199
	distance	*	-13.947	1.029	-13.554	0.000	-16.015	-11.982
	size	*	3.363	0.280	12.025	0.000	2.816	3.913
	distance * size	*	31.443	3.041	10.341	0.000	25.570	37.493
<b>Faster R-CNN</b>								
Vanishing	object	*	-1.040	0.113	-9.179	0.000	-1.266	-0.821
	distance	*	-27.692	2.080	-13.316	0.000	-31.877	-23.729
	size	*	2.888	0.384	7.528	0.000	2.142	3.646
	distance * size		-13.106	7.590	-1.727	0.084	-28.136	1.621
Mislabeling	object	*	-1.009	0.141	-7.175	0.000	-1.292	-0.740
	distance	*	-23.256	2.228	-10.440	0.000	-27.775	-19.048

Untargeted	size	*	1.152	0.411	2.805	0.005	0.348	1.958
	distance * size		3.841	7.680	0.500	0.617	-11.396	18.717
	object	*	-0.530	0.090	-5.900	0.000	-0.708	-0.356
	distance	*	-18.563	1.238	-14.997	0.000	-21.049	-16.197
	size	*	2.349	0.285	8.254	0.000	1.793	2.909
<hr/>								
<b>Cascade R-CNN</b>								
Vanishing	object	*	-1.019	0.100	-10.149	0.000	-1.218	-0.824
	distance	*	-33.425	2.225	-15.020	0.000	-37.885	-29.166
	size	*	5.056	0.422	11.968	0.000	4.241	5.898
	distance * size	*	-40.938	8.862	-4.619	0.000	-58.577	-23.819
Mislabeling	object	*	-0.898	0.123	-7.302	0.000	-1.143	-0.660
	distance	*	-31.929	2.726	-11.713	0.000	-37.433	-26.754
	size	*	2.875	0.392	7.341	0.000	2.111	3.646
	distance * size		-10.216	9.136	-1.118	0.263	-28.219	7.593
Untargeted	object	*	-0.642	0.097	-6.605	0.000	-0.835	-0.454
	distance	*	-29.381	1.881	-15.623	0.000	-33.163	-25.791
	size	*	1.493	0.306	4.884	0.000	0.895	2.094
	distance * size	*	32.873	5.009	6.563	0.000	23.086	42.736

---

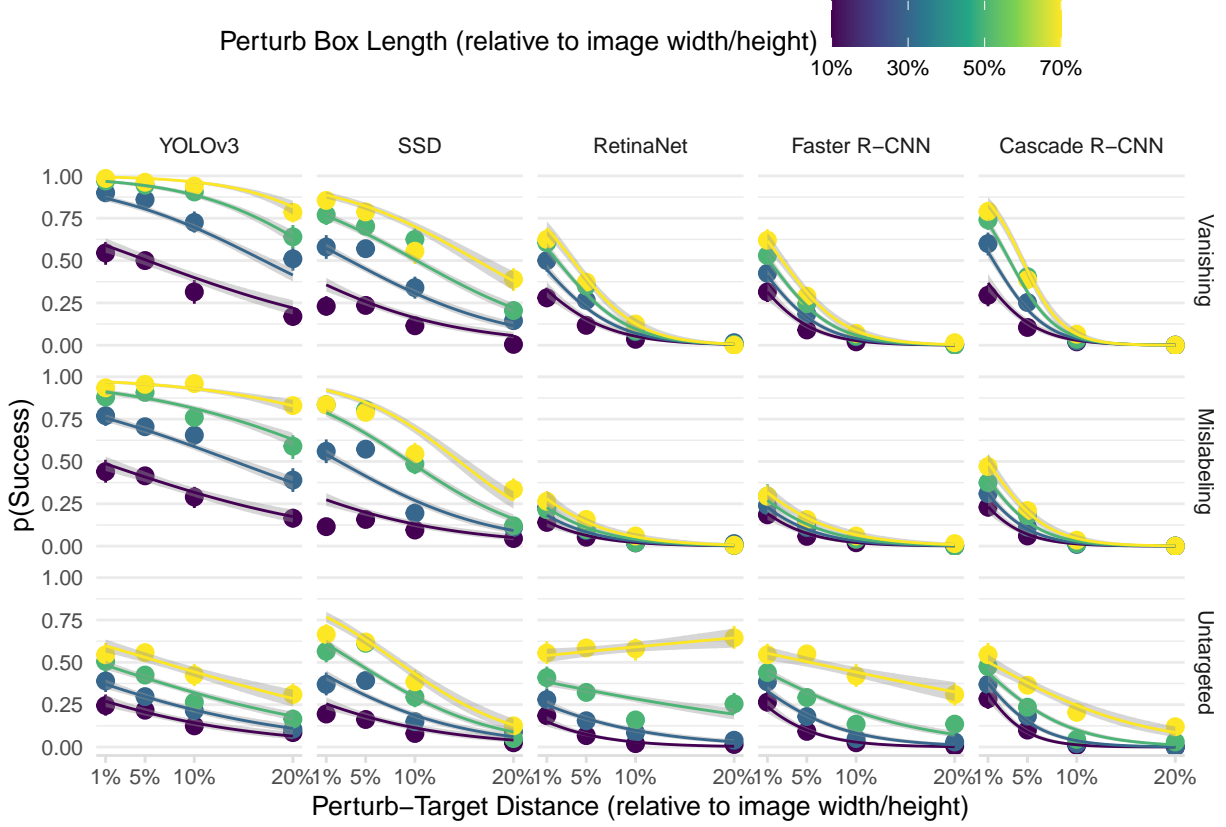


Figure 1: **A deliberate attack obfuscates intent with increased success for all models and attacks:** We implement intent obfuscating attack by perturbing an arbitrary non-overlapping square region to disrupt a randomly selected target object at various lengths and distances. The binned summaries and regression trendlines graph success proportion against perturb-target distance (relative to image width/height) and perturb box length (relative to image width/height) in the deliberate attack experiment. Errors are 95% confidence intervals, and every point aggregates success over 200 images. The deliberate attack multiplies success as compared to the randomized attack (Figure ??), especially at close perturb-target distance (relative to image width/height) and large perturb box length (relative to image width/height). Full details are given in Section ??.